

Vehicle Tutorial Chapter 1: Getting Started

Ekaterina Komendantskaya and Matthew Daggitt (today's presentors), on behalf of the Vehicle team



We will:

 ... introduce Main building blocks of Vehicle as a Programming Language



We will:

- ... introduce Main building blocks of Vehicle as a Programming Language
- ... get to practice working with Vehicle



We will:

- ... introduce Main building blocks of Vehicle as a Programming Language
- ▶ ... get to practice working with Vehicle
- ... use the famous ACAS Xu benchmark to show Vehicle's work flow – from specification to verification



We will:

- ... introduce Main building blocks of Vehicle as a Programming Language
- ▶ ... get to practice working with Vehicle
- ... use the famous ACAS Xu benchmark to show Vehicle's work flow – from specification to verification
- ... try to identify PL problems that are being resolved by Vehicle

Recap: four PL problems



Recap: four PL problems



- I^O Interoperability properties are not portable between training/counter-example search/ verification.
- I^P Interpretability code is not easy to understand.
- I^{\int} Integration properties of networks cannot be linked to larger control system properties.
- E^G Embedding gap little support for translation between problem space (as in original spec) and input space (at neural network level).

Recap: ACAS Xu



A collision avoidance system for unmanned autonomous aircraft.

Inputs:

- \triangleright Distance to intruder, ρ
- ightharpoonup Angle to intruder, θ
- ▶ Intruder heading, φ
- ightharpoonup Speed, v_{own}
- ▶ Intruder speed, v_{int}

Outputs:

- ► Clear of conflict
- ► Strong left
- ► Weak left
- ► Weak right
- ► Strong right

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

10 different specified properties in total.

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

10 different specified properties in total.

Definition (ACAS Xu: Property 1)

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

10 different specified properties in total.

Definition (ACAS Xu: Property 1)

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

$$(\rho \ge 55947.691) \land (v_{own} \ge 1145) \land (v_{int} \le 60)$$

 \Rightarrow the score for COC is at most 1500

Table of Contents



Vehicle' Syntax

Types

Let us build the ACAS Xu specification. We start with types of input and output vectors, as well as types of ACAS Xu networks

```
type InputVector = Vector Rat 5
type OutputVector = Vector Rat 5
```

@network

acasXu : InputVector -> OutputVector

The Vector type represents a mathematical vector, or in programming terms can be thought of as a fixed-length array.

Values



Types for values are automatically inferred by **Vehicle**. For example, we can declare the number π and its type will be inferred as rational:

pi = 3.141592

Working with vectors

some input or output pre-processing maybe expected when defining a neural network.

Example

It is assumed that the ACAS Xu inputs and outputs are normalised, i.e. the network does not work directly with units like m/s. However, the specifications we want to write should ideally concern the original units.

Working with vectors

some input or output pre-processing maybe expected when defining a neural network.

Example

It is assumed that the ACAS Xu inputs and outputs are normalised, i.e. the network does not work directly with units like m/s. However, the specifications we want to write should ideally concern the original units.

- ► This is an instance of "problem space / input space mismatch"
- ▶ ... that is very common in neural net verification
- ▶ Being able to reason about problem space (alongside the input space) is a feature that distinguishes **Vehicle** from majority of the mainstream neural network verifiers

Vector normalisation

For clarity, we define a new type synonym for unnormalised input vectors which are in the problem space.

```
type UnnormalisedInputVector = Vector Rat 5
```

Next we define the range of the inputs that the network is designed to work over.

```
minimumInputValues : UnnormalisedInputVector
minimumInputValues = [0,0,0,0,0]
```

```
\verb|maximumInputValues|: UnnormalisedInputVector|
```

maximumInputValues = [60261.0, 2*pi, 2*pi, 1100.0, 1200.0]

```
{\tt meanScalingValues} \; : \; {\tt UnnormalisedInputVector}
```

meanScalingValues = [19791.091, 0.0, 0.0, 650.0, 600.0]

Vector manipulation

An alternative method to vector definition is to use the 'foreach' constructor, which is used to provide a value for each 'index i'.

```
\label{limit_minimum} \begin{tabular}{ll} minimumInputValues : UnnormalisedInputVector \\ minimumInputValues = for each i . 0 \end{tabular}
```

Let us see how 'foreach' works with vector indexing. We can now define the normalisation function that takes an input vector and returns the unnormalised version.

```
normalise : UnnormalisedInputVector -> InputVector
normalise x = foreach i .
  (x ! i - meanScalingValues ! i) / (maximumInputValues ! i)
```

Vector manipulation

An alternative method to vector definition is to use the 'foreach' constructor, which is used to provide a value for each 'index i'.

```
\label{limit_minimum} \begin{tabular}{ll} minimumInputValues : UnnormalisedInputVector \\ minimumInputValues = for each i . 0 \end{tabular}
```

Let us see how 'foreach' works with vector indexing. We can now define the normalisation function that takes an input vector and returns the unnormalised version.

```
normalise : UnnormalisedInputVector -> InputVector
normalise x = foreach i .
   (x ! i - meanScalingValues ! i) / (maximumInputValues ! i)
... our first acquaintance with functions!
```



```
<name> : <type> <name> [<args>] = <expr>
```

Functions make up the backbone of the **Vehicle** language.



```
<name> : <type> <name> [<args>] = <expr>
```

Functions make up the backbone of the **Vehicle** language.

```
validInput : UnnormalisedInputVector -> Bool
validInput x = forall i .
  minimumInputValues ! i <= x ! i <= maximumInputValues ! i</pre>
```



```
<name> : <type> <name> [<args>] = <expr>
```

Functions make up the backbone of the **Vehicle** language.

```
validInput : UnnormalisedInputVector -> Bool
validInput x = forall i .
  minimumInputValues ! i <= x ! i <= maximumInputValues ! i</pre>
```

Our first acquaintance with quantifiers!

One of the main advantages of **Vehicle** is that it can be used to state and prove specifications that describe the network's behaviour over an infinite set of values.



Function Composition: Exercise!

Infer the types of functions 'acasXu' and 'normalise':

normAcasXu : UnnormalisedInputVector -> OutputVector

normAcasXu x = acasXu (normalise x)

Pre-defined functions and predicates



We have already used:

!

<=

Exercise

What do they stand for?

Lets verify ACAS Xu!

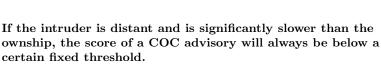


```
distanceToIntruder = 0 -- measured in metres
angleToIntruder = 1 -- measured in radians
intruderHeading = 2 -- measured in radians
speed = 3 -- measured in metres/second
intruderSpeed = 4 -- measured in meters/second

clearOfConflict = 0
weakLeft = 1
weakRight = 2
strongLeft = 3
strongRight = 4
```

The fact that all vector types come annotated with their size means that it is impossible to mess up indexing into vectors, e.g. if you changed 'distanceToIntruder = 0' to 'distanceToIntruder = 5' the specification would fail to type-check.

Property 1





Property 1

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

```
intruderDistantAndSlower : UnnormalisedInputVector -> Bool
intruderDistantAndSlower x =
   x ! distanceToIntruder >= 55947.691 and
   x ! speed >= 1145 and
   x ! intruderSpeed <= 60</pre>
```

Property 1

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

```
intruderDistantAndSlower : UnnormalisedInputVector -> Bool
intruderDistantAndSlower x =
   x ! distanceToIntruder >= 55947.691 and
   x ! speed >= 1145 and
   x ! intruderSpeed <= 60</pre>
```

Exercise!

- 1. Can you identify whether the specification is written in terms of input space or problem space? How do you know?
- 2. Can you spot more pre-defined **Vehicle** functions? What are they?

There is little left to do!

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.



There is little left to do!

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

There is little left to do!

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

Exercise!

1. Can you guess the purpose of the syntax

```
@property
```

?

2. What kind of domain 'forall' ranges over? Is it finite or infinite?

How to run Vehicle



Checklist

- 1. a verifier installed (Marabou);
- 2. the actual network is supplied in an ONNX format
- 3. Vehicle is installed.

How to run Vehicle



Checklist

- 1. a verifier installed (Marabou);
- 2. the actual network is supplied in an ONNX format
- 3. Vehicle is installed.

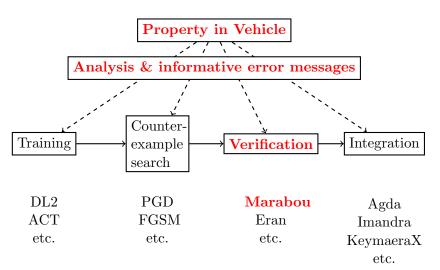
Exercise: ϵ -ball Robustness

```
type Image = Tensor Rat [28, 28]
type Label = Index 10
validImage : Image -> Bool
validImage x = forall i i . 0 <= x ! i ! i <= 1
@network
classifier : Image -> Vector Rat 10
advises : Image -> Label -> Bool
advises x i = forall j . j != i => classifier x ! i > classifier x ! j
@parameter
epsilon : Rat
boundedByEpsilon : Image -> Bool
boundedByEpsilon x = forall i j . -epsilon <= x ! i ! j <= epsilon
robustAround : Image -> Label -> Bool
robustAround image label = forall pertubation .
 let perturbedImage = image - pertubation in
 boundedByEpsilon pertubation and validImage perturbedImage =>
    advises perturbedImage label
@dataset
trainingImages : Vector Image n
@dataset
trainingLabels : Vector Label n
@property
robust : Vector Bool n
robust = foreach i . robustAround (trainingImages ! i) (trainingLabels ! i)
```



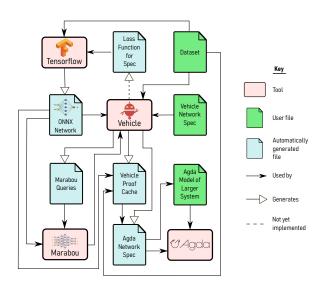
Vehicle ...

is a domain-specific functional language for writing high-level property specifications for neural networks



Vehicle Architecture





Sources





M. Daggitt, R. Atkey, W. Kokke, E. Komendantskaya, L. Arnaboldi: Compiling Higher-Order Specifications to SMT Solvers: How to Deal with Rejection Constructively. CPP 2023



N. Slusarz, E. Komendantskaya, M. Daggitt, R. Stewart, K. Stark: Logic of Differentiable Logics: Towards a Uniform Semantics of DL. LPAR 2023.



Matthew L. Daggitt, Wen Kokke, Robert Atkey, Luca Arnaboldi, Ekaterina Komendantskaya: Vehicle: Interfacing Neural Network Verifiers with Interactive Theorem Provers. FOMLAS



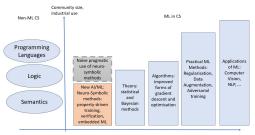
Vehicle Team: The Vehicle language: https://github.com/vehicle-lang 2023.



M.Daggitt and W.Kokke: Vehicle User Manual. 2023.

Which challenges Vehicle addresses

- ▶ Theory: finding appropriate verification properties
- Solvers: undecidability of non-linear real arithmetic and scalability of neural network verifiers
- ▶ ML: understanding and integrating property-driven training
- ▶ Programming: finding the right languages to support these developments
- ► Complex systems: integration of neural net verification into complex systems



ML methods, increasing applied nature