

Vehicle Tutorial Chapter 1: Getting Started

Ekaterina Komendantskaya and Matthew Daggitt and Wen Kokke and ...

Vehicle Team



Matthew Daggitt





Bob Atkey



Rob Stewart



Kathrin Stark



Marco Casadio







Luca Arnaboldi



Table of Contents



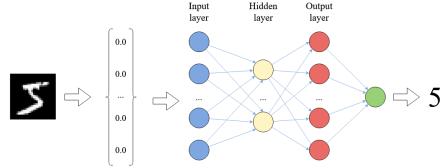
Neural Network Verification: overview of the new domain

The lifecycle of neural network verification

Introduction to Vehicle (a mini tutorial)

Neural nets for classification





Formally,

a neural network is a function $N: \mathbb{R}^n \to \mathbb{R}^m$.



... are ideal for "perception" tasks:

- ▶ approximate functions when exact solution is hard to get
- ▶ tolerant to noisy and incomplete data



... are ideal for "perception" tasks:

- ▶ approximate functions when exact solution is hard to get
- ▶ tolerant to noisy and incomplete data



... are ideal for "perception" tasks:

- ▶ approximate functions when exact solution is hard to get
- ▶ tolerant to noisy and incomplete data



... are ideal for "perception" tasks:

- ▶ approximate functions when exact solution is hard to get
- ▶ tolerant to noisy and incomplete data

BUT

- ➤ solutions not easily conceptualised (lack of explainability)
- ▶ prone to a new range of safety and security problems:



... are ideal for "perception" tasks:

- ▶ approximate functions when exact solution is hard to get
- ▶ tolerant to noisy and incomplete data

BUT

- ▶ solutions not easily conceptualised (lack of explainability)
- ▶ prone to a new range of safety and security problems: adversarial attacks data poisoning catastrophic forgetting



























the perturbations are imperceptible to human eye









the perturbations are imperceptible to human eye attacks transfer from one neural network to another





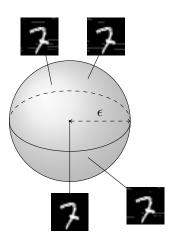




the perturbations are imperceptible to human eye attacks transfer from one neural network to another affect any domain where neural networks are applied

Verification Property: " ϵ -ball robustness"





An
$$\epsilon$$
-ball $\mathbb{B}(\hat{\mathbf{x}}, \epsilon) = {\mathbf{x} \in \mathbb{R}^n : |\hat{\mathbf{x}} - \mathbf{x}| \le \epsilon}$

Classify all points in $\mathbb{B}(\hat{\mathbf{x}}, \epsilon)$ "robustly".

Another example property: ACAS Xu



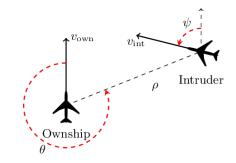
A collision avoidance system for unmanned autonomous aircraft.

Inputs:

- \triangleright Distance to intruder, ρ
- \triangleright Angle to intruder, θ
- \blacktriangleright Intruder heading, φ
- ightharpoonup Speed, v_{own}
- ightharpoonup Intruder speed, v_{int}

Outputs:

- ► Clear of conflict
- ► Strong left
- ► Weak left
- ► Weak right
- Strong right



ACAS Xu

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

ACAS Xu

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

10 different specified properties in total.

ACAS Xu

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

10 different specified properties in total.

Definition (ACAS Xu: Property 1)

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

ACAS X11

The system was originally implemented as a 2Gb lookup table but was replaced with a neural network in order to improve size and latency requirements.

10 different specified properties in total.

Definition (ACAS Xu: Property 1)

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

$$(\rho \ge 55947.691) \land (v_{own} \ge 1145) \land (v_{int} \le 60)$$

 \Rightarrow the score for COC is at most 1500

More Generally



Given $N: \mathbb{R}^n \to \mathbb{R}^m$

Verification of such functions most commonly boils down to specifying admissible intervals for the function's output given an interval for its inputs.

More Generally



Given $N: \mathbb{R}^n \to \mathbb{R}^m$

Verification of such functions most commonly boils down to specifying admissible intervals for the function's output given an interval for its inputs.



Casadio, M., Komendantskaya, E., Daggitt, M.L., Kokke, W., Katz, G., Amir, G., Refaeli, I.: Neural network robustness as a verification property: A principled case study. In: Computer Aided Verification (CAV 2022).

Overview of The Verification Landscape





Panda Gibbon Paper Szegedy et al. (2013)

I have this specification I want to verify!





What tools are available? 2015





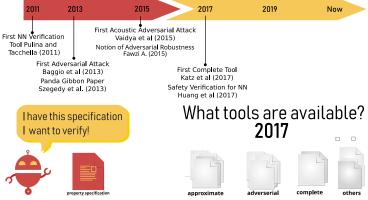


complete

others

Overview of The Verification Landscape





Overview of The Verification Landscape





\$

A whole range of domain-specific verifiers exist:

A whole range of domain-specific verifiers exist:

► Marabou (SMT technology)



A whole range of domain-specific verifiers exist:

- ► Marabou (SMT technology)
- ► ERAN (abstract interpretation + MILP)





A whole range of domain-specific verifiers exist:

- ► Marabou (SMT technology)
- ► ERAN (abstract interpretation + MILP)
- ► Verisig (interval arithmetic)
- ► AlphaBetaCROWN (linear bound propagation)
- **.** . . .

International Standards and Competitions

https://www.vnnlib.org/

\$

A whole range of domain-specific verifiers exist:

- ► Marabou (SMT technology)
- ► ERAN (abstract interpretation + MILP)
- ► Verisig (interval arithmetic)
- ► AlphaBetaCROWN (linear bound propagation)
- **.** . . .

International Standards and Competitions

https://www.vnnlib.org/

Marabou is our choice as it is complete, and the set of expressible queries is large!



Guy Katz, Clarke Barrett, D. Dill, K. Julian, and M. Kochenderfer. Reluplex: An Efficient SMT Solver for Verifying Deep Neural Networks. In CAV, 2017.

Table of Contents



Neural Network Verification: overview of the new domain

The lifecycle of neural network verification

Introduction to Vehicle (a mini tutorial)



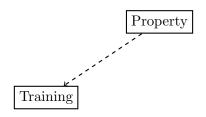
Property



Property

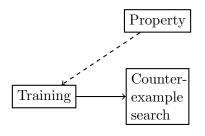
Training





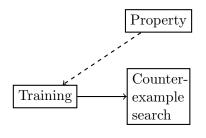
DL2 ACT etc.





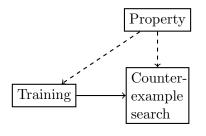
DL2 ACT etc.





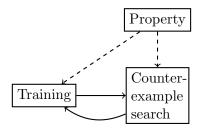
DL2 ACT etc.





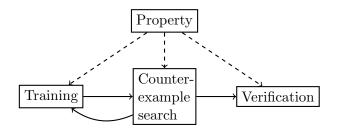
DL2 PGD ACT FGSM etc. etc.





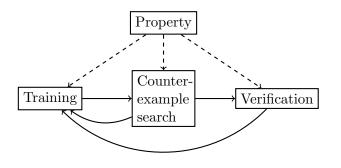
DL2 PGD ACT FGSM etc. etc.





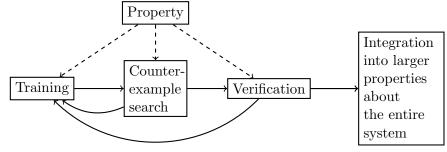
DL2 ACT etc. PGD FGSM etc. Marabou Eran etc.





DL2 ACT etc. PGD FGSM etc. Marabou Eran etc.





DL2 ACT etc. PGD FGSM etc. Marabou Eran etc.



▶ Theory: finding appropriate verification properties



- ▶ Theory: finding appropriate verification properties
- ➤ Solvers: undecidability of non-linear real arithmetic and scalability of neural network verifiers



- ▶ Theory: finding appropriate verification properties
- ➤ Solvers: undecidability of non-linear real arithmetic and scalability of neural network verifiers
- ► ML: understanding and integrating property-driven training



- ▶ Theory: finding appropriate verification properties
- ▶ Solvers: undecidability of non-linear real arithmetic and scalability of neural network verifiers
- ► ML: understanding and integrating property-driven training
- ▶ Programming: finding the right languages to support these developments



- ► Theory: finding appropriate verification properties
- ➤ Solvers: undecidability of non-linear real arithmetic and scalability of neural network verifiers
- ► ML: understanding and integrating property-driven training
- Programming: finding the right languages to support these developments
- Complex systems: integration of neural net verification into complex systems



How BIG is the problem?

Lets look under the hood...

Training framework: DL2

```
$
```

```
class RobustnessConstraint(Constraint):
126
127
141
         def get domains(self, x batches, v batches):
              assert len(x batches) == 1
142
143
              n_batch = x_batches[0].size()[0]
144
145
              return [[Box(np.clip(x batches[0][i].cpu().numpv() - self.eps, 0, 1),
146
                           np.clip(x batches[0][i].cpu().numpy() + self.eps, 0, 1))
147
                      for i in range(n batch)]]
148
149
         def get_condition(self, z_inp, z_out, x_batches, y_batches):
150
              n_batch = x_batches[0].size()[0]
151
              z_out = transform_network_output(z_out, self.network_output)[0]
152
              #z logits = F.log softmax(z out[0], dim=1)
153
154
              pred = z out[np.arange(n batch), v batches[0]]
155
156
              limit = torch.FloatTensor([0.3])
157
             if self.use cuda:
158
                  limit = limit.cuda()
159
              return dl2.GEO(pred, torch.log(limit))
```



Fischer, M., Balunovic, M., Drachsler-Cohen, D., Gehr, T., Zhang, C., and Vechev, M. T. DL2: training and querying neural networks with logic. In Proc. of the 36th Int. Conf. Machine Learning, ICML 2019

Training framework: ART



```
@classmethod
333
334
          def property6a(cls, dom: AbsDom):
              p = AcasProp(name='property6a', dom=dom, safe fn='cols is min', viol fn='cols not min',
                           fn args=[AcasOut.CLEAR OF CONFLICT])
336
337
              p.set input bound(AcasIn.RHO, new low=12000, new high=62000)
             p.set input bound(AcasIn.THETA, new low=0.7, new high=3.141592)
338
339
              p.set input bound(AcasIn.PSI, new low=-3.141592, new high=-3.141592 + 0.005)
              p.set input bound(AcasIn.V OWN, new low=100, new high=1200)
340
              p.set input bound(AcasIn.V INT, new low=0, new high=1200)
341
342
             p.set_all_applicable_as(False)
343
              p.set_applicable(1, 1, True)
344
              return p
```



Lin, X., Zhu, H., Samanta, R., and Jagannathan, S. (2020). Art: Abstraction refinement-guided training for provably correct neural networks. In FMCAD 2020

Verification framework: Marabou



```
def test_acas_1_normalize():
    """

Test the 1,1 experimental ACAS Xu network.

By passing "normalize=true" to read_nnet, Marabou adjusts the parameters of the first and last layers of the network to incorporate the normalization.

As a result, properties can be defined in the original input/output spaces without any manual normalization.

"""

filename = "acasxu/ACASXU_experimental_v2a_1_1.nnet"
testInputs = [
    [1800.0, 0.0, -1.5, 180.0, 180.0],
    [19080.0, -3.0, -1.5, 380.0, 380.0],
    [58080.0, -3.0, 0.0, 380.0, 680.0]]
]

testOutputs = [
    [177.87553729, 173.75796115, 193.85928806, 153.87876146, 195.80495022],
    [-8.55188079, 0.48863711, 0.44258938, 0.44315198, 0.43950133],
    [29.9190734, 27.2386958, 45.82497222, 14.5619455, 46.86448056]
]
network = evaluateFile(filename, testInputs, testOutputs, normalize = True)
```



Katz, G., Huang, D. A., Ibeling, D., Julian, K., Lazarus, C., Lim, R., Shah, P., Thakoor, S., Wu, H., Zeljic, A., Dill, D. L., Kochenderfer, M. J., and Barrett, C. W. (2019). The Marabou framework for verification and analysis of deep neural networks. In CAV 2019

Verification framework: ERAN



```
1 [12000, 62000]
2 [0.7, 3.141592][-3.141592, -0.7]
3 [-3.141592, -3.136592]
4 [100, 1200]
5 [0, 600]
1 5
2 y0 min
```



Singh, G., Gehr, T., Püschel, M., and Vechev, M. T. (2019). An abstract domain for certifying neural networks. PACMPL, 3(POPL):41:1–41:30.

Verification property language: VNNLIB



```
28
     (assert (or
29
         (and (<= X 0 0.700434925) (>= X 0 -0.129289109)
30
              (<= X 1 0.499999896) (>= X 1 0.11140846)
31
              (<= X 2 -0.499204121) |(>= X 2 -0.499999896)
32
              (\le X \ 3 \ 0.5) \ (\ge X \ 3 \ -0.5)
33
              (<= X 4 0.5) (>= X 4 -0.5))
34
         (and (<= X 0 0.700434925) (>= X 0 -0.129289109)
35
              (<= X 1 -0.11140846) (>= X 1 -0.499999896)
36
              (<= X 2 -0.499204121) (>= X 2 -0.499999896)
37
              (<= X_3 0.5) (>= X_3 -0.5)
38
              (\le X 4 0.5) (\ge X 4 - 0.5))
39
     ))
40
41
     : unsafe if coc is not minimal
42
     (assert (or
43
         (and (<= Y 1 Y 0))
44
        (and (<= Y 2 Y 0))
45
         (and (<= Y 3 Y 0))
46
         (and (<= Y 4 Y 0))
47
    ))
48
```





► Interoperability - properties are not portable between training/counter-example search/verification.



- ► Interoperability properties are not portable between training/counter-example search/verification.
- ► Interpretability code is not easy to understand.



- ► Interoperability properties are not portable between training/counter-example search/verification.
- ► Interpretability code is not easy to understand.
- ► Integration properties of networks cannot be linked to larger control system properties.

Vehicle ...

is a domain-specific functional language for writing high-level property specifications for neural networks

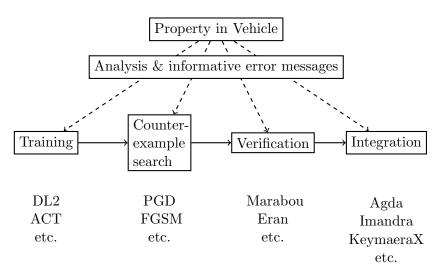


Table of Contents



Neural Network Verification: overview of the new domain

The lifecycle of neural network verification

Introduction to Vehicle (a mini tutorial)

Types

Let us build the ACAS Xu specification. We start with types of input and output vectors, as well as types of ACAS Xu networks

```
type InputVector = Vector Rat 5
type OutputVector = Vector Rat 5
```

@network

acasXu : InputVector -> OutputVector

The Vector type represents a mathematical vector, or in programming terms can be thought of as a fixed-length array.

Values



Types for values are automatically inferred by **Vehicle**. For example, we can declare the number π and its type will be inferred as rational:

pi = 3.141592

Working with Vectors

some input or output pre-processing maybe expected when defining a neural network.

Example

It is assumed that the ACAS Xu inputs and outputs are normalised, i.e. the network does not work directly with units like m/s. However, the specifications we want to write should ideally concern the original units.

Working with Vectors

some input or output pre-processing maybe expected when defining a neural network.

Example

It is assumed that the ACAS Xu inputs and outputs are normalised, i.e. the network does not work directly with units like m/s. However, the specifications we want to write should ideally concern the original units.

- ► This is an instance of "problem space / input space mismatch"
- ▶ ... that is very common in neural net verification
- ▶ Being able to reason about problem space (alongside the input space) is a feature that distinguishes **Vehicle** from majority of the mainstream neural network verifiers

Vector Normalisation

For clarity, we define a new type synonym for unnormalised input vectors which are in the problem space.

```
type UnnormalisedInputVector = Vector Rat 5
```

Next we define the range of the inputs that the network is designed to work over.

```
minimumInputValues : UnnormalisedInputVector
minimumInputValues = [0,0,0,0,0]
```

```
\verb|maximumInputValues|: UnnormalisedInputVector|
```

maximumInputValues = [60261.0, 2*pi, 2*pi, 1100.0, 1200.0]

```
{\tt meanScalingValues} \; : \; {\tt UnnormalisedInputVector}
```

meanScalingValues = [19791.091, 0.0, 0.0, 650.0, 600.0]

Vector manipulation

An alternative method to vector definition is to use the 'foreach' constructor, which is used to provide a value for each 'index i'.

```
\label{limit_minimum} \begin{tabular}{ll} minimumInputValues : UnnormalisedInputVector \\ minimumInputValues = for each i . 0 \end{tabular}
```

Let us see how 'foreach' works with vector indexing. We can now define the normalisation function that takes an input vector and returns the unnormalised version.

```
normalise : UnnormalisedInputVector -> InputVector
normalise x = foreach i .
  (x ! i - meanScalingValues ! i) / (maximumInputValues ! i)
```

Vector manipulation

An alternative method to vector definition is to use the 'foreach' constructor, which is used to provide a value for each 'index i'.

```
\label{limit_minimum} \begin{tabular}{ll} minimumInputValues : UnnormalisedInputVector \\ minimumInputValues = for each i . 0 \end{tabular}
```

Let us see how 'foreach' works with vector indexing. We can now define the normalisation function that takes an input vector and returns the unnormalised version.

```
normalise : UnnormalisedInputVector -> InputVector
normalise x = foreach i .
   (x ! i - meanScalingValues ! i) / (maximumInputValues ! i)
... our first acquaintance with functions!
```



```
<name> : <type> <name> [<args>] = <expr>
```

Functions make up the backbone of the **Vehicle** language.



```
<name> : <type> <name> [<args>] = <expr>
```

Functions make up the backbone of the **Vehicle** language.

```
validInput : UnnormalisedInputVector -> Bool
validInput x = forall i .
  minimumInputValues ! i <= x ! i <= maximumInputValues ! i</pre>
```



```
<name> : <type> <name> [<args>] = <expr>
```

Functions make up the backbone of the **Vehicle** language.

```
validInput : UnnormalisedInputVector -> Bool
validInput x = forall i .
  minimumInputValues ! i <= x ! i <= maximumInputValues ! i</pre>
```

Our first acquaintance with quantifiers!

One of the main advantages of **Vehicle** is that it can be used to state and prove specifications that describe the network's behaviour over an infinite set of values.



Function Composition: Exercise!

normAcasXu : UnnormalisedInputVector -> OutputVector

normAcasXu x = acasXu (normalise x)

Pre-defined functions and predicates



We have already used:

!

<=

Exercise

What do they stand for?

Lets verify ACAS Xu!



```
distanceToIntruder = 0 -- measured in metres
angleToIntruder = 1 -- measured in radians
intruderHeading = 2 -- measured in radians
speed = 3 -- measured in metres/second
intruderSpeed = 4 -- measured in meters/second

clearOfConflict = 0
weakLeft = 1
weakRight = 2
strongLeft = 3
strongRight = 4
```

The fact that all vector types come annotated with their size means that it is impossible to mess up indexing into vectors, e.g. if you changed 'distanceToIntruder = 0' to 'distanceToIntruder = 5' the specification would fail to type-check.

Property 1

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.



Property 1

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

```
intruderDistantAndSlower : UnnormalisedInputVector -> Bool
intruderDistantAndSlower x =
   x ! distanceToIntruder >= 55947.691 and
   x ! speed >= 1145 and
   x ! intruderSpeed <= 60</pre>
```

Property 1

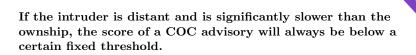
If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

```
intruderDistantAndSlower : UnnormalisedInputVector -> Bool
intruderDistantAndSlower x =
   x ! distanceToIntruder >= 55947.691 and
   x ! speed >= 1145 and
   x ! intruderSpeed <= 60</pre>
```

Exercise!

- 1. Can you identify whether the specification is written in terms of input space or problem space? How do you know?
- 2. Can you spot more pre-defined **Vehicle** functions? What are they?

There is little left to do!



There is little left to do!

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

There is little left to do!

If the intruder is distant and is significantly slower than the ownship, the score of a COC advisory will always be below a certain fixed threshold.

Exercise!

1. Can you guess the purpose of the syntax

```
@property
?
```

2. What kind of domain 'forall' ranges over? Is it finite or infinite?

How to run Vehicle



Checklist

- 1. a verifier installed (Marabou);
- 2. the actual network is supplied in an ONNX format
- 3. Vehicle is installed.

How to run Vehicle



Checklist

- 1. a verifier installed (Marabou);
- 2. the actual network is supplied in an ONNX format
- 3. Vehicle is installed.

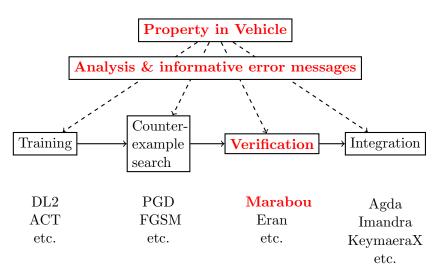
Exercise: ϵ -ball Robustness

```
type Image = Tensor Rat [28, 28]
type Label = Index 10
validImage : Image -> Bool
validImage x = forall i i . 0 <= x ! i ! i <= 1
@network
classifier : Image -> Vector Rat 10
advises : Image -> Label -> Bool
advises x i = forall j . j != i => classifier x ! i > classifier x ! j
@parameter
epsilon : Rat
boundedByEpsilon : Image -> Bool
boundedByEpsilon x = forall i j . -epsilon <= x ! i ! j <= epsilon
robustAround : Image -> Label -> Bool
robustAround image label = forall pertubation .
 let perturbedImage = image - pertubation in
 boundedByEpsilon pertubation and validImage perturbedImage =>
    advises perturbedImage label
@dataset
trainingImages : Vector Image n
@dataset
trainingLabels : Vector Label n
@property
robust : Vector Bool n
robust = foreach i . robustAround (trainingImages ! i) (trainingLabels ! i)
```



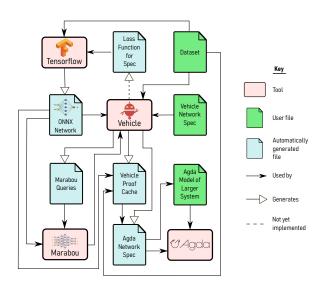
Vehicle ...

is a domain-specific functional language for writing high-level property specifications for neural networks



Vehicle Architecture





Sources





M. Daggitt, R. Atkey, W. Kokke, E. Komendantskaya, L. Arnaboldi: Compiling Higher-Order Specifications to SMT Solvers: How to Deal with Rejection Constructively. CPP 2023



N. Slusarz, E. Komendantskaya, M. Daggitt, R. Stewart, K. Stark: Logic of Differentiable Logics: Towards a Uniform Semantics of DL. LPAR 2023.



Matthew L. Daggitt, Wen Kokke, Robert Atkey, Luca Arnaboldi, Ekaterina Komendantskaya: Vehicle: Interfacing Neural Network Verifiers with Interactive Theorem Provers. FOMLAS



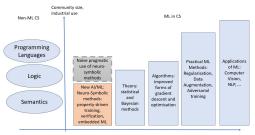
Vehicle Team: The Vehicle language: https://github.com/vehicle-lang 2023.



M.Daggitt and W.Kokke: Vehicle User Manual. 2023.

Which challenges Vehicle addresses

- ▶ Theory: finding appropriate verification properties
- Solvers: undecidability of non-linear real arithmetic and scalability of neural network verifiers
- ▶ ML: understanding and integrating property-driven training
- ▶ Programming: finding the right languages to support these developments
- ► Complex systems: integration of neural net verification into complex systems



ML methods, increasing applied nature